
Task 5: Write a query using an aggregate function.

Query:

```
hive> select max(id) as max_id, min(id) as min_id from College;
```

RunProof:

```
hive> select max(id) as max_id, min(id) as min_id from College;
Query ID = anagave_20210503231722_d480c5c0-f943-4354-a631-5c1d554bb9c4
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1620056543801_0074, Tracking URL = http://hadoop-
nn001.cs.okstate.edu:8088/proxy/application\_1620056543801\_0074/
Kill Command = /usr/local/hadoop-3.3.0/bin/mapred job -kill
job_1620056543801_0074
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 1
2021-05-03 23:17:30,002 Stage-1 map = 0%, reduce = 0%
2021-05-03 23:17:35,286 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 2.48 sec
2021-05-03 23:17:36,343 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 21.24
sec
2021-05-03 23:17:41,596 Stage-1 map = 100%, reduce = 100%, Cumulative CPU
25.57 sec
MapReduce Total cumulative CPU time: 25 seconds 570 msec
Ended Job = job_1620056543801_0074
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 1 Cumulative CPU: 25.57 sec HDFS Read: 45894
HDFS Write: 107 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 570 msec
OK
106 101
Time taken: 20.267 seconds, Fetched: 1 row(s)
```

Output:

```
Total MapReduce CPU Time Spent: 25 seconds 570 msec
OK
106 101
Time taken: 20.267 seconds, Fetched: 1 row(s)
```