

Task 6: Write a query using a nested SELECT statement

Query:

```
hive> select * from College where id in( select id from mascot where mascot = "Pistol Pete");
```

RunProof:

```
Query ID = anagave_20210503231947_9c6bb0e4-57c3-4462-a07c-8ee59c0ad23c
Total jobs = 1
SLF4J: Found binding in [jar:file:/usr/local/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/
org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFactory]2021-05-03 23:19:55 Starting to launch local task to process map join;
maximum memory = 4750049280
2021-05-03 23:19:58 End of local task; Time Taken: 2.191 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1620056543801_0075, Tracking URL = http://hadoop-
nn001.cs.okstate.edu:8088/proxy/application_1620056543801_0075/
Kill Command = /usr/local/hadoop-3.3.0/bin/mapred job -kill job_1620056543801_0075
Hadoop job information for Stage-3: number of mappers: 5; number of reducers: 0
2021-05-03 23:20:07,174 Stage-3 map = 0%, reduce = 0%
2021-05-03 23:20:18,720 Stage-3 map = 60%, reduce = 0%, Cumulative CPU 12.85 sec
2021-05-03 23:20:19,772 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 20.62 sec
MapReduce Total cumulative CPU time: 20 seconds 620 msec
Ended Job = job_1620056543801_0075
MapReduce Jobs Launched:
Stage-Stage-3: Map: 5 Cumulative CPU: 20.62 sec HDFS Read: 53852 HDFS Write: 511
SUCCESS
Total MapReduce CPU Time Spent: 20 seconds 620 msec
OK
101 Kim [2127893564,2121123546] {"city":"Stillwater","state":"Oklahoma"}
{"COA":2,"OS":2} US
Time taken: 33.826 seconds, Fetched: 1 row(s)
```

Output:

```
Total MapReduce CPU Time Spent: 20 seconds 620 msec
OK
101 Kim [2127893564,2121123546] {"city":"Stillwater","state":"Oklahoma"} {"COA":2,"OS":2} US
Time taken: 33.826 seconds, Fetched: 1 row(s)
```